LESSON 7 · MID-TERM · PRODUCT PROTOTYPE

Callback

We help you own the room — with clear, provable words.

For STEM students first: turn real work into evidence you can explain, alone, in the interview.

Fictional customer: Mei Lin · Financial figures are labeled assumptions · Callback does not predict hiring

Strong results, weak narration

BEFORE · MEI

“I’m a data science student. I know Python, SQL, and Power BI.”

True — and it proves nothing. No one used it, nothing changed, no link to the job.

✓

The work is real

Projects, internships, and research with results the student did produce.

?

Follow-ups expose the gap

Who used it? What changed? What did you decide? Answers fall apart.

!

Generic AI writes for you

Polished answers that overclaim — and that you can’t defend live.

Results are the proof. Narration is the missing skill.

One loop from resume to a defensible answer

1 Paste resume + JD

Read on your device; contact details can be removed first.

2 Pick up to 3 job needs

Each need quotes the exact JD sentence.

3 Map your evidence

Only facts you supplied and confirmed. Qualifiers stay.

4 Answer + up to 3 follow-ups

Probes what's missing: who used it, what changed, fit.

5 Feedback that quotes you

Four rubric scores, one priority fix, your own words.

6 Retry, coaching hidden

Compare first vs. retry; diagnosis says what to practice.

Three things that make it different

AI follow-ups that probe, not praise

Answer one question. Get up to 3 follow-ups on what's missing — who used it, what changed, how it fits this role.

"Who used the dashboard, and what did it help them monitor?"

Language support for international students

Think it through in your preferred language first. Nothing is added to your story until you confirm it.

Write details in Chinese → confirm the fact → practice in English.

Feedback as rewrite + progress tracking

Not just a score. A quoted priority fix, a retry with coaching hidden, and first vs. retry side by side.

Mei's relevance: first 2 → retry 4 (fictional prepared example).

Use case: Mei Lin applies to Bayline Freight

FICTIONAL CUSTOMER

Mei Lin

M.S. Data Science, Bayview University (expected 2028)

Target: Data Analyst Intern, Summer 2027

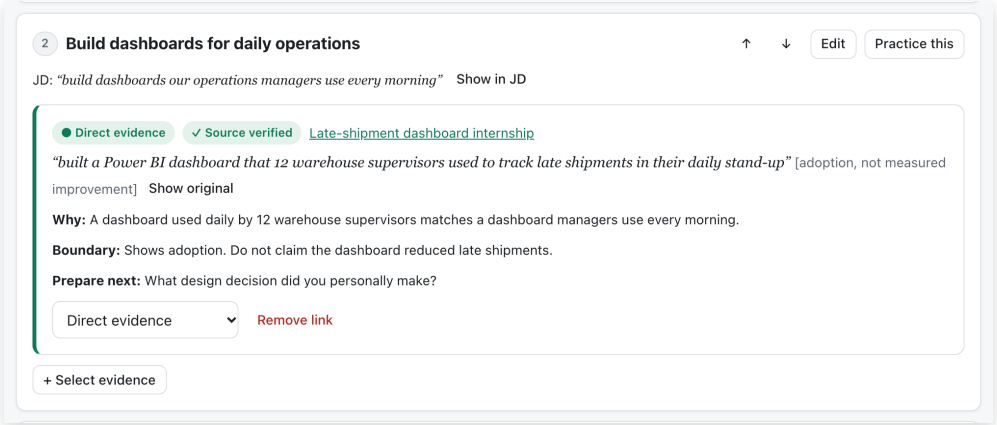
20 minutes a day to prepare

JD need

“build dashboards our operations managers use every morning”

Her evidence

Power BI dashboard used by 12 warehouse supervisors in their daily stand-up.



Value map in the prototype: the JD quote and her resume quote are both verified against the source. “Direct evidence” — with the boundary written down.

Mei’s drill: probe, retry, compare

Interviewer
Tell me about a time your work supported a team’s daily operational decisions.

Mei
During my internship, I built a Power BI dashboard for the warehouse team.

Follow-up 1 · evidence
Who used the dashboard, and what did it help them monitor?

Follow-up 2 · outcome
What can you say about the outcome of that work?

Follow-up 3 · relevance
How does that experience connect to this internship?

Criterion	First	Retry
Relevance	2	4
Evidence and ownership	2	3
Organization	3	4
Clear language	3	4

Same question, coaching hidden. Prepared fictional scores.

INDEPENDENT RETRY

“…a Power BI dashboard that 12 warehouse supervisors used to track late shipments in their daily stand-up. … I cannot claim it reduced delivery delays.”

Market fit: who it’s for and what we add

WHO

STEM students preparing for internship and first-job interviews

International students first: real results, interviews in a second language.

WHO PAYS

Students (low-cost plan). Career centers later, after pilots.

What students need	Generic AI chat	Mock-interview apps	Callback
Claims you can defend	Writes freely, can overclaim	Varies by product	**Every claim cites a confirmed fact**
Follow-ups on what’s missing	Only if you ask	Often offered	**Up to 3, targeted**
Proof you can say it alone	Not tracked	Partly	**Hidden-coaching retry + comparison**
Think in your own language	Yes	Varies	**Yes — confirmed before use**

Category-level view from public product pages; not tested across every plan.

Company goal and pilot objectives

Goal: every STEM student can explain their real work clearly and honestly — without a script.

O1 · Prove learning

In blind counselor review, the retry is rated clearer than the first answer in most pairs.

Target: > 60% of pairs

O2 · Prove demand

Run campus pilots and see students finish the full loop, including the unassisted retry.

Target: 2 pilots, 100 students, 40% finish a retry

O3 · Prove the model

Keep AI cost per active student low enough for a student-priced plan.

Target: AI cost < $2.50 per student per month

Targets for the pilot phase, not results. No customers or measured outcomes yet.

Roadmap aligned to the goals

Phase 1 — now
STEM student interviews

Serves O1 · O2 · O3

Phase 2 — later
All majors

Grows O2 once O1 is proven

Phase 3 — later
Career advancement, promotion reviews

New segment after O3 holds

Built in the Phase 1 prototype

- Value map with source-verified quotes
- Drill: ≤ 3 follow-ups, hidden-coaching retry
- Rubric feedback + first-vs-retry comparison
- Practice diagnosis, team-priority simulation
- Prep-time planner within your minutes
- AI backend with output validation

Next, before the first pilot

- Verify live AI on real student answers (O1)
- Speech-to-text for spoken practice (O1)
- Accounts + opt-in sharing with counselors (O2)
- Cost tracking per student (O3)

Pricing: free pilots, then a student plan

Campus pilot

Free

One term with a cohort. We cover AI costs to measure O1 and O2.

Student plan

$9 / month

Price hypothesis to test. Practice drills, retries, and progress tracking.

Career center license

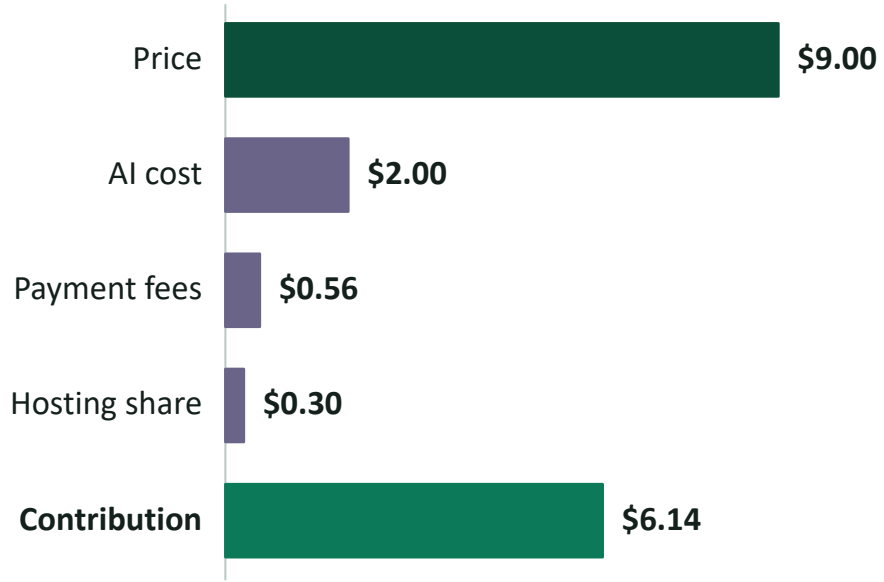
Not decided

An option to test after pilots, not decided.

How we test the price in pilots: a willingness-to-pay survey and the share of pilot students who choose the paid plan. No customers or prices yet — both are hypotheses.

Unit economics per paying student (assumptions)

Monthly, per paying student (USD)



$9.00 − $2.00 − $0.56 − $0.30 = $6.14 contribution

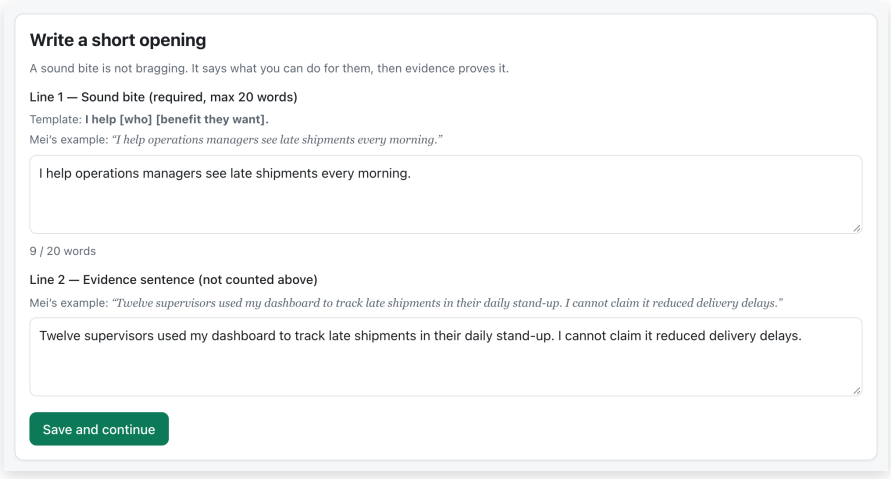
≈ 82

paying students a month to cover $500 of fixed costs at $6.14 contribution each

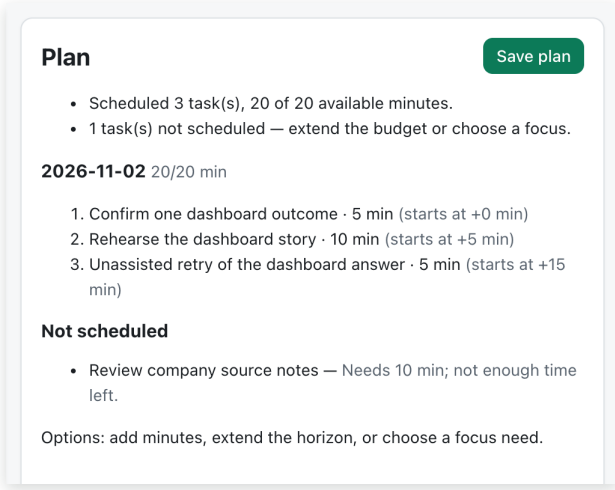
Assumptions

- AI: Claude Opus 5.5 list price, $4 / $20 per million input / output tokens
- One drill ≈ 13.5k input + 6.4k output tokens ≈ $0.18
- 8 drills + other AI help a month ≈ $2.00
- Card fees 2.9% + $0.30; team unpaid in pilots

Prototype: what the live demo shows



Opening: a sound bite (≤ 20 words), then evidence.



Planner: 20 minutes → A+B+C scheduled, D left out.

Demo path: Home → Try Mei's example → Value map → Drill with follow-ups → First vs. retry → Plan my prep time.

Built with Claude Code from a 17-step prompt kit: React + TypeScript, local-first storage, Claude Opus 5.5 backend with validated output, 42 automated tests.

THE ASK

Help us run the first pilot

One campus career center for a free spring pilot, and ten STEM students to try the drill this month.

What Callback does not do

No answers during a real interview. Preparation only.

No hiring predictions.

Never adds achievements you did not supply and confirm.

We help you own the room — with clear, provable words.